

## **1. Introduction:**

For this midterm, the goal was to develop a machine learning model to predict star ratings on a scale from 1.0 to 5.0 of Amazon movie reviews on both textual and metadata attributes. Using the provided kaggle dataset, we aimed to train a model off of 1,697,533 unique reviews, and ultimately fill in and identify the scores of reviews where the final rating was missing.

The included kaggle dataset included the following attributes:

- ProductId - unique identifier for the product
- UserId - unique identifier for the user
- HelpfulnessNumerator - number of users who found the review helpful
- HelpfulnessDenominator - number of users who indicated whether they found the review helpful
- Score - rating between 1 and 5
- Time - timestamp for the review
- Summary - brief summary of the review
- Text - text of the review
- Id - a unique identifier associated with a review


Ultimately, I made note that there are 2 general categories of the data: non-textual; Floats and Integers that we can use with less processing, and textual data, Text and Summary attributes that make up the bulk of the dataset. I attempted to experiment around with the non-textual data first. This is because I knew that manipulating the textual data for use in the model would be difficult and be a costly operation, given the size of the dataset. Hence, I experimented, and found the following features to be most important, before returning to the textual data and doing the more heavy analysis later.

## **2. Data and methodology**

### **a. Features**


The majority of the time I spent on developing this machine learning model was in developing features; and then a smaller subset of the time was spent working with the parameters and hyperparameters of the end model I decided to use.

Firstly, regarding the features. I created several simple features that require no explanation: Length of review, time of month and year of the review. Furthermore, I decided to add an average product score attribute based on each unique UserId's given score. In cases where there were no scores to attribute, I populated with the average score of the dataset: **4.1105**.



Distribution of scores of the dataset


Finally of note is Helpfulness as an attribute; I used *HelpfulnessNumerator*, *HelpfulnessDenominator*, and finally an aggregate helpfulness attribute of the numerator over the denominator, which I called *Helpfulness*.



Now in terms of the text themselves, I focused on two things: Sentiment analysis, and vectorization. Vectorization was done as covered in class; running `tf_idf` to obtain 500 vectors of the most common words, and calculating the average occurrences of each word throughout the dataset. Sentiment analysis was done through the `TextBlob` library, which is a library built on python's `NLTK` library. It parses through each Text and assigns a given score of sentiment and polarity based on a database of keywords. Positive keywords raise the sentiment, negative keywords decrease the sentiment, and a negation word(not good, rarely impressive) reverses the sentiment of the following keyword. Together, these represented the bulk of the features that had a final impact on the outcome of the model.

## b. Model

After testing, I decided to use a Random Forest Classifier.



I ultimately tested Logistic Regression models, as well as K-Nearest Neighbors classifiers, but none performed as well as RFC. RFC is better because of the amount of features used; all 500 of the vectorized words are present. In KNN this amount of features would cause it to not be able to compile the model; and in logistic regression it would cause too much noise. I believe that the Decision trees of Random Forest Classifiers are perfect for this use case.

### c. Conclusion

Ultimately, I obtained an accuracy on the testing set of **0.629489**. Through some rudimentary hyper parameter tuning, I used 300 estimators of the model, over 510 features.

### 3. Discussion and conclusion.

I ultimately believe that the biggest constraint of the project was time. There was a lot more I could have done in terms of both the features in the model. In future projects, I will look to run TF\_IDF on more samples; I could have run on both the summary and the text itself. I also would have liked to run more statistical analysis on the scores and the non-textual data; I could have found correlations between more attributes of the data, and created more creative features of the data.

I didn't have the time I would have liked to to do proper hyper parameter testing. I would've run RandomizedSearchCV to iterate through different variations of the model, expanding out of different scoring systems, max features, estimators, and more.

With more time I would have also expanded more on the model itself; I would have probably used an aggregate of different models to more accurately assess parts of the data.